

First 90 Days

Senior Product Manager, Agentic AI | Russell Dudek

Map the behavior. Prove the thin slice. Create the compounding loop.

A hypothesis-driven entry plan. Exact priorities, baselines, and targets must be confirmed with Jerry's product, engineering, operations, licensed experts, and leadership.

Days 0-30: See the real system

WORKSTREAM	ACTIONS	EVIDENCE / OUTPUT
Behavior inventory	Map the six prompt locations to customer intents, owners, models, tools, data, handoffs, evaluation, release paths, and known failure classes.	Behavior/source map; ownership map; change-path diagram; list of unknowns and duplicated policy.
Customer and operational evidence	Review conversations, escalations, licensed-agent handoffs, repeat contacts, tool failures, and support/operations feedback by intent.	Failure taxonomy linked to customer consequence, operational burden, and learning opportunity.
Baseline scorecard	Separate quality, automation, escalation, latency, cost, and CSAT by intent and surface; identify where averages hide tradeoffs.	Intent-level baseline with data-quality caveats and decision-ready definitions.
Thin-slice selection	Choose one high-volume, bounded journey with meaningful customer value, manageable authority, clear tools, and measurable outcome.	Signed problem statement, customer promise, decision rights, and pilot success/stop conditions.
Stakeholder compact	Clarify working cadence among product, engineering, applied AI, operations, licensed experts, analytics, legal/compliance, and leadership.	Decision-rights matrix and escalation route; no committee theater.

Questions answered by day 30

- Where does behavior drift originate, and where is it merely observed?
- Which failure modes are customer-visible, silently operational, or economically expensive?
- What does “automated” mean at Jerry, and what quality debt can hide inside that metric?
- Which authority is model-owned, tool-owned, licensed-human-owned, or leadership-owned?
- Which existing assets should be preserved rather than replaced?

Location workstream
Complete Tampa Bay relocation logistics and establish Florida residency on a candid, mutually confirmed timeline.

From proof to platform

Days 31-90

Days 31-60: Install the manifest and ship one bounded behavior

ACTION	DEFINITION OF DONE
Define the minimum Agent Behavior Manifest	Customer promise, intent/context, tool authority, human authority, evaluation suite, cost/latency budget, owner/release policy, and learning path are explicit and versioned.
Build the evaluation path	Representative cases, edge conditions, tool failures, handoff behavior, shadow/canary observation, and post-release monitoring are connected to release decisions.
Clarify release posture	Ship, canary, hold, or human-first decisions have evidence requirements, accountable owners, and rollback criteria.
Release the thin slice	One journey improves against a balanced scorecard without hiding a decline in customer trust, economics, latency, or operational burden.
Run the learning review	Observed failures are classified; fixes and product decisions are linked; reusable platform requirements are separated from one-off exceptions.

Days 61-90: Turn proof into a compounding platform

ACTION	DEFINITION OF DONE
Separate core behavior from surface adapters	The team can explain what is shared across chat, voice, computer use, vision, and back-office workflows - and what is legitimately surface-specific.
Prioritize roadmap bets	Voice, ChatGPT, computer use, and platform consolidation are compared using customer value, technical readiness, risk, economics, and reusable learning.
Establish product standards	Prompt/context changes, model upgrades, tool permissions, evaluation, ownership, and rollback follow a lightweight but explicit standard.
Create the operating cadence	Weekly evidence review, release decisions, incident learning, and monthly platform/roadmap decisions have clear participants and inputs.
Define next-quarter outcomes	Targets are grounded in confirmed baselines and specify quality, customer outcome, automation integrity, economics, change safety, and reuse.

What I would not do

No rewrite-first strategy Preserve working leverage and prove the behavioral abstraction before large-scale migration.	No single vanity metric Automation rate is important, but not at the expense of false completion, customer trust, or hidden human work.	No capability theater Voice and computer use become roadmap commitments only when a customer journey, authority model, evidence path, and economics are credible.
---	--	--

Day-90 leadership readout

A concise, decision-ready account of the current behavior system, the thin-slice result, customer and operational learnings, platform standards adopted, unresolved risks, relocation status, and the recommended next set of product bets.