

Agent Behavior Manifest

A versioned product definition for one customer-facing agent behavior

One customer promise. One governed behavior. Many deliberate surfaces.

Use this worksheet before a prompt or model change becomes a production behavior. It is intentionally decision-oriented: unknowns should remain visible rather than being disguised as maturity scores.

Behavior identity

Behavior name / version	_____
Customer promise	What outcome is the agent responsible for? What must it never imply it completed?
Primary intent	Which customer need activates this behavior? What neighboring intents are explicitly out of scope?
Journey state / prerequisites	Required data, consent, authentication, policy state, or prior actions.
Surfaces	Chat / voice / ChatGPT / computer use / vision / back office / other: _____
Accountable product owner	_____ Technical owner: _____

Authority and action

Permitted tools / APIs	Actions, inputs, validation, confirmation, and permission boundaries.
Prohibited actions	Actions the agent may not perform or claim to perform.
Human authority	Where a licensed expert, operator, reviewer, or leader decides, approves, or owns the outcome.
Handoff trigger	Conditions, urgency, customer language, required context, and ownership after handoff.
Failure behavior	How the system responds to unavailable tools, ambiguous state, stale data, policy conflict, or partial completion.
Recovery / retry	Retry limit, alternate path, customer communication, and incident escalation.

Evidence before release

EVIDENCE CLASS	REQUIRED CASES / DECISION	OWNER
Core task cases	Representative customer intents, successful tool use, truthful completion, and expected response/action.	_____
Edge / ambiguity	Conflicting data, incomplete information, neighboring intent, unsupported request, and policy exception.	_____
Human handoff	Correct trigger, sufficient context, no duplicated questioning, and clear ownership transfer.	_____
Tool failure	Timeout, unavailable API, invalid response, retry behavior, and customer-safe recovery.	_____
Adversarial / misuse	Prompt injection, data leakage, prohibited action, social engineering, and misleading completion.	_____

Release and learning

Evidence, economics, ownership, and change control

Release scorecard

DIMENSION	BASELINE / THRESHOLD	OBSERVED EVIDENCE	DECISION OWNER
Answer / action quality	_____	_____	_____
Customer resolution / CSAT	_____	_____	_____
Automation integrity	_____	_____	_____
Escalation quality	_____	_____	_____
Latency	_____	_____	_____
Model + tool cost / resolved intent	_____	_____	_____
Operational burden	_____	_____	_____

Release posture

SHIP Coverage, authority, economics, monitoring, owner, and rollback are sufficient for broad release.	CANARY Bounded exposure with explicit observation window, stop conditions, and rollback.	HOLD / HUMAN-FIRST Required evidence or authority is missing; limits are stated and safe human ownership is preserved.
Selected posture / rationale	_____ _____	
Exposure / canary cohort	_____	
Stop conditions	_____	
Rollback owner / method	_____	

Production learning loop

Signals reviewed	Conversation outcomes, tool traces, escalations, complaints, repeat contacts, licensed-agent feedback, operations/support incidents, cost and latency.
Failure classification	Intent / context / policy / tool / model / prompt / data / handoff / interface / ownership / measurement.
Learning review cadence	Daily operational exceptions: _____ Weekly evidence review: _____ Monthly product decision: _____
Change record	Observed evidence → decision → manifest version → evaluation change → release → production result.
Reusable learning	What becomes a platform standard, channel adapter, evaluation fixture, tool policy, or human-authority rule?

Approval record

ROLE	NAME	DECISION / DATE
Product		
Engineering / applied AI		
Operations / licensed authority		
Risk / compliance as required		