

Interview Thesis Brief

Senior Product Manager, Agentic AI | Russell Dudek

Scale the behavior, not the prompt.

Jerry's real mandate is to convert an early, homegrown AI advantage into a versioned customer product platform that can improve quality, automation, and customer satisfaction while controlling latency, cost, operational risk, and ownership complexity.

Public operating scale >70% of inbound sales and service automated; more than 50,000 chats per month.	Platform inflection Prompt logic across six locations with no single source of truth or clear home for voice, computer use, and consumer AI bets.	Business context 5M+ customers, profitability, broadening from auto insurance into physical-asset ownership.
---	---	--

The real mandate

Own the product system that defines how Jerry's agents should behave, proves that behavior before release, adapts it across surfaces, and learns from production. This role is simultaneously platform PM, customer-experience architect, AI evaluation leader, and technical translator.

Central contradictions

TENSION	FAILURE IF MISHANDLED	PRODUCT RESPONSE
Consolidation vs. velocity	A platform rewrite slows active teams or preserves fragmentation indefinitely.	Start with one high-value thin slice; standardize the behavior object before centralizing every implementation.
Autonomy vs. human authority	The agent appears capable but mishandles exceptions, regulated judgment, or customer trust.	Encode tool limits, licensed authority, handoff context, and failure behavior as first-class product requirements.
Quality vs. cost and latency	Model upgrades improve one dimension while damaging economics or experience.	Use a balanced release scorecard and intent-level budgets rather than a single blended quality score.
Shared standard vs. surface fit	Voice, chat, vision, and computer use either drift or become forced into one interaction pattern.	Separate a shared behavioral core from channel-specific adapters and evaluation suites.

The operating model

Agent Behavior Manifest

A versioned product object for agent behavior

The manifest is not another prompt repository. It is the minimum product definition required to release a customer-facing behavior responsibly and repeatedly.

Customer promise	Outcome the agent owns, including what it must never imply it completed.
Intent and context	Customer need, journey state, data prerequisites, and activation conditions.
Tool authority	Permitted APIs/actions, preconditions, confirmation rules, retries, and failure behavior.
Human authority	Where licensed experts or operators decide, approve, recover, or assume ownership.
Evaluation suite	Offline cases, edge conditions, shadow evidence, canary thresholds, and production monitors.
Cost / latency budget	Economics and response-time constraints by intent and surface.
Owner / release policy	Accountable owner, version, rollout posture, rollback, and incident response.
Learning loop	How failures become classified evidence, decisions, validated changes, and reusable knowledge.

Release postures, not false precision

Ship Behavior is covered, economics are understood, authority is explicit, and rollback is ready.	Canary Promising change with bounded exposure, defined observation, and a fast rollback path.	Hold / human-first Required evidence, authority, or recovery behavior is missing; the system states limits honestly.
---	---	--

Balanced product scorecard

DIMENSION	DECISION QUESTIONS
Answer and action quality	Did the agent interpret the intent, use the right tools, follow policy, and complete the task truthfully?
Customer outcome	Was the issue resolved, was escalation useful, and did the customer need to contact Jerry again?
Automation integrity	Which automated completions are real, which escalations are avoidable, and where is failure silent?
Economics	What is the latency, model/tool cost, and human-review cost per resolved intent?
Change safety	What regression evidence, canary evidence, rollback, and incident learning support the release?
Platform reuse	Can another surface adopt the behavior without duplicating policy, evaluation, or ownership logic?

Interview agenda

Discovery questions, evidence, and proof tests

Executive discovery questions

1. Where does fragmented prompt ownership currently create the most expensive failure: customer inconsistency, release risk, engineering rework, economics, or roadmap delay?
2. Which automation metric is most likely to look healthy while hiding a customer or operational failure mode?
3. Which journey offers the best thin-slice proof for a unified behavior standard without becoming a rewrite?
4. How are licensed authority, escalation context, and post-handoff learning represented today?
5. What must be true for voice or computer use to become a roadmap commitment rather than a model-capability experiment?
6. What evidence would make the COO say this role has become a material growth lever?

Evidence I would bring to the discussion

EVIDENCE BASE	RELEVANT OPERATING PATTERN
DudeWorth	AI-readiness, agentic workflows, retrieval, human review, escalation, evaluation, governance, and rapid capability experimentation.
Vape-Jet	Cross-system ownership, telemetry, quality, customer/support/engineering feedback, and adoption across approximately 1,000 fielded robots.
Amazon	Customer-backward mechanisms, standard work, escalation, and reliable high-velocity execution.
Compunetics	Consequence-aware engineering, quality systems, root cause, technical translation, and controls.

Strongest objection and proof test

Objection: no conventional startup PM or forward-deployed engineering title sequence. **Response:** do not accept a title-equivalence argument. Test whether I can define one complete agent behavior, trace dependencies and failure modes, formulate the SQL and evaluation questions, make quality/cost/latency tradeoffs legible, and align engineering with the operators and licensed experts who inherit the outcome.

De-risking commitment

30-day evidence map. 60-day thin-slice release. 90-day roadmap and operating cadence grounded in actual customer, technical, economic, and human-authority evidence.

Location

Currently in Pittsburgh and actively relocating to Tampa Bay. Prepared to establish Florida residency and confirm timing directly; no claim of current Florida residency is made.